

The ish Shell

Implement the Unix shell described in the accompanying man page. Submit your solution (source files only) in a tarball on moodle. Be sure to include all the files needed to make the shell. Your shell should compile by typing `make`, producing an executable file named `ish`. You will be penalized if your code contains unresolvable external references.

To help you get started, I suggest that you become familiar with the Unix system calls defined in Section 2 of the Unix programmer's manual. You should also be aware that there are several library routines (Section 3) that offer convenient interfaces to some of the more cryptic system calls. However, **you may not use the library routine `system()`, nor any of the routines prohibited in the `ish` man page**. Chapters 8, 9, and parts of 10 in the Stevens book are also very relevant.

You may, if you desire, use `lex` and `yacc` to implement your command parsing. Even if you have not used these programs before, you will probably find it much simpler to learn them than to write your own parser. The `lex` and `yacc` input files `ish.l` and `ish.y`, respectively, are provided, along with tutorials.

To implement `ish` you will obviously be creating a process that forks off other processes, which may in turn fork processes, etc. Unix limits the number of processes you can have; if you run out you won't be able to do very much so be careful. While you are testing you may want to somehow limit the number of shells you can have running at the same time.

When in doubt about the functionality of `ish`, or how it should behave in a particular situation, model its behavior after that of `csh`. Questions of general interest should be posted to moodle.

NAME

ish – a shell (command interpreter) with a csh-like syntax and advanced interactive features.

SYNOPSIS

ish

DESCRIPTION

ish, a shell, is a command interpreter with syntax similar to csh.

INITIALIZATION AND TERMINATION

When first started, ish performs commands from the file `~/ishrc`, provided that it is readable. Typically, the `~/ishrc` file contains commands to specify the terminal type and environment.

INTERACTIVE OPERATION

After startup processing, an interactive ish shell begins reading commands from the terminal, prompting with *hostname%*. The shell then repeatedly performs the following actions: a line of command input is read and broken into words; this sequence of words is parsed (as described under USAGE); and the shell executes each command in the current line.

USAGE**LEXICAL STRUCTURE**

The shell splits input lines into words separated by spaces or tabs, with the following exceptions:

- * The special characters '&', '|', '<', sequences of special characters form single words: '>>',
- * Special characters preceded by a backslash '\ ' character prevents the shell from interpreting them as special characters.
- * Strings enclosed in double quotes (") or quotes (') form part or all of a single word. Special characters inside of strings do not form separate words.

COMMAND LINE PARSING

A simple command is a sequence of words, the first of which specifies the command to be executed. A pipeline is a sequence of one or more simple commands separated by | or |&. With '|', the standard output of the preceding command is redirected to the standard input of the command that follows. With '|&', both the standard error and the standard output are redirected through the pipeline. A list is a sequence of one or more pipelines separated by ; or &. These separators have the following meanings:

- ;
Causes sequential execution of the preceding pipeline. The shell waits for the pipeline to finish. A newline character following a pipeline behaves the same as a ';'.
- &
Causes asynchronous execution of the preceding pipeline. The shell does not wait for the pipeline to finish; instead, it displays the job number (see Job Control) and associated process IDs, and begins processing the next pipeline (prompting if necessary).

I/O REDIRECTION

The following metacharacters indicate that the subsequent word is the name of a file to which the standard input, standard output, and/or standard error is redirected.

- <
Redirect the standard input.
- >, >&
Redirect the standard output to a file. If the file does not exist, it is created. If it does exist, it is overwritten; its previous contents are lost. The '&' form redirects both standard output and the standard error to the file.
- >>, >>&
Append the standard output. Like '>', but places output at the end of the file rather than overwriting it. The '&' form appends both the standard error and standard output to the file.

COMMAND EXECUTION

If the command is an ish shell built-in, the shell executes it directly. Otherwise, the shell searches for a file by that name with execute access. If the command-name contains a '/', the shell takes it as a pathname, and searches for it. If the command-name does not contain a '/', the shell attempts to resolve it to a pathname, searching each directory in the PATH variable for the command.

When a file is found that has proper execute permissions, the shell forks a new process and passes it, along with its arguments, to the OS using the `execve` system call (you must use this system call). The OS then attempts to overlay the new process with the desired program. If the file is an executable binary the OS succeeds, and begins executing the new process.

If the file does not have execute permissions, or if the pathname matches a directory, a 'permission denied' message is displayed. If the pathname cannot be resolved a 'command not found' message is displayed. If either of these errors occur with any component of a pipeline the entire pipeline is aborted, although some of the components of the pipeline may have already started running.

A pipeline is completed (i.e. returns to the prompt), only when all the commands that form a part of the pipeline and that are being executed in the foreground are completed.

ENVIRONMENT VARIABLES

Environment variables may be accessed via the `setenv` and `unsetenv` built-in commands. Initially no environment variables are set, i.e. ish does not inherit environment variables from its parent. Ish maintains environment variables internally, and must not use the C library routines `putenv` or `getenv`. When a program is exec'ed the environment variables are passed as parameters to `execve`. The only environment variable that ish interprets is `PATH`; all other environment variables can be set and unset in ish using the above builtin commands, but are not interpreted.

SIGNAL HANDLING

The shell normally ignores `QUIT` signals. Background jobs are immune to signals generated from the keyboard, including hangups (`HUP`). Other signals have the values that ish inherited from its environment. Shells catch the `TERM` signal.

JOB CONTROL

The shell associates a numbered job with each command sequence, to keep track of those commands that are running in the background or have been stopped with `TSTP` signals (typically `CTRL-Z`). Jobs are put into the foreground using the `tcsetpgrp` system call. When a command is started in the background using the `'&'` metacharacter, the shell displays a line with the job number in brackets, and a list of associated process numbers; e.g.,

```
[1] 1234
```

To see the current list of jobs, use the `jobs` built-in command. The job most recently stopped (or put into the background if none are stopped) is referred to as the current job.

To manipulate jobs, refer to the built-in commands `bg`, `fg`, and `kill`.

A reference to a job begins with a `'%'`. Refer to job number `j` as in: `'kill %j'`.

A job running in the background stops when it attempts to read from the terminal. Background jobs can normally produce output, but this can be suppressed using the `'stty tostop'` command.

STATUS REPORTING

While running interactively, the shell tracks the status of each job and reports whenever it finishes or becomes blocked. Status only need be reported prior to printing the prompt; this means that job status may be polled via `wait`. Do not use signal handlers to track job status.

BUILT-IN COMMANDS

Built-in commands are executed within *ish*. If a built-in command occurs as any component of a pipeline except the last, it is executed in a subshell.

bg %job ...

Run the specified job in the background.

cd [dir]

Change the shell's working directory to directory dir. If no argument is given, change to the home directory of the user.

exit

Causes ish to exit. ✓

fg %job

Bring the specified job into the foreground.

jobs

List the active jobs under job control. →

kill %job ...

Send the TERM (terminate) signal to the job indicated. To insure termination, the job is also sent a CONT (continue) signal.

✓ **setenv [VAR [word]]**

With no arguments, setenv displays all environment variables. With the VAR argument, it sets the environment variable VAR to have an empty (null) value. (By convention, environment variables are normally given upper-case names.) With both VAR and word arguments, setenv sets the named environment variable to the value word, which must be either a single word or a quoted string.

✓ **unsetenv VAR**

Remove VAR from the environment.

FILES

~/.ishrc Read at beginning of execution by each shell.